

Untitled

l

Installing package into 'C:/Users/vasud/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)

package 'ggthemes' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\vasud\AppData\Local\Temp\RtmpMFYmLY\downloaded_packages

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --

v dplyr 1.1.4 v readr 2.1.5

v forcats 1.0.0 v stringr 1.5.1

v ggplot2 3.5.1 v tibble 3.2.1

v lubridate 1.9.4 v tidyr 1.3.1

v purrr 1.0.2

-- Conflicts ----- tidyverse_conflicts() --

x dplyr::filter() masks stats::filter()

x dplyr::lag() masks stats::lag()

i Use the conflicted package (<<http://conflicted.r-lib.org/>>) to force all conflicts to become errors

Attaching package: 'data.table'

The following objects are masked from 'package:lubridate':

hour, isoweek, mday, minute, month, quarter, second, wday, week,
yday, year

The following objects are masked from 'package:dplyr':

between, first, last

The following object is masked from 'package:purrr':

transpose

Attaching package: 'janitor'

The following objects are masked from 'package:stats':

chisq.test, fisher.test

Attaching package: 'jsonlite'

The following object is masked from 'package:purrr':

flatten

Attaching package: 'plotly'

The following object is masked from 'package:httr':

config

The following object is masked from 'package:ggplot2':

last_plot

The following object is masked from 'package:stats':

filter

The following object is masked from 'package:graphics':

layout

corrplot 0.95 loaded

Loading required package: zoo

Attaching package: 'zoo'

The following objects are masked from 'package:data.table':

yearmon, yearqtr

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:httr':

progress

The following object is masked from 'package:purrr':

lift

Attaching package: 'MASS'

The following object is masked from 'package:plotly':

select

The following object is masked from 'package:dplyr':

select

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Registered S3 method overwritten by 'quantmod':

method from
as.zoo.data.frame zoo

Attaching package: 'forecast'

The following object is masked from 'package:ggpubr':

gghistogram

Loading required package: car

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:dplyr':

recode

The following object is masked from 'package:purrr':

some

Loading required package: survival

Attaching package: 'survival'

The following object is masked from 'package:caret':

cluster

Loading required package: StanHeaders

rstan version 2.32.6 (Stan version 2.32.2)

For execution on a local, multicore CPU with excess RAM we recommend calling
options(mc.cores = parallel::detectCores()).
To avoid recompilation of unchanged Stan programs, we recommend calling
rstan_options(auto_write = TRUE)
For within-chain threading using `reduce_sum()` or `map_rect()` Stan functions,
change `threads_per_chain` option:
rstan_options(threads_per_chain = 1)

Do not specify '-march=native' in 'LOCAL_CPPFLAGS' or a Makevars file

Attaching package: 'rstan'

The following object is masked from 'package:tidyr':

extract

Loading required package: Rcpp

Loading 'brms' package (version 2.22.0). Useful instructions
can be found by typing help('brms'). A more detailed introduction
to the package is available through vignette('brms_overview').

Attaching package: 'brms'

The following object is masked from 'package:rstan':

loo

The following object is masked from 'package:survival':

kidney

The following object is masked from 'package:forecast':

ma

The following object is masked from 'package:lme4':

ngrps

The following object is masked from 'package:stats':

ar

Attaching package: 'coda'

The following object is masked from 'package:rstan':

traceplot

Welcome to BayesFactor 0.9.12-4.7. If you have questions, please contact Richard Morey (richarddmorey@gmail.com)

Type BFManual() to open the manual.

Loading required package: antitrust

Attaching package: 'antitrust'

The following object is masked from 'package:car':

logit

The following object is masked from 'package:forecast':

CV

Attaching package: 'trade'

The following object is masked from 'package:antitrust':

sim

Attaching package: 'kableExtra'

The following object is masked from 'package:dplyr':

group_rows

Attaching package: 'officer'

The following object is masked from 'package:readxl':

read_xlsx

Attaching package: 'xgboost'

The following object is masked from 'package:plotly':

slice

The following object is masked from 'package:dplyr':

slice

Loaded glmnet 4.1-8

-- Attaching packages ----- tidymodels 1.2.0 --

v broom	1.0.7	v rsample	1.2.1
v dials	1.3.0	v tune	1.2.1

```

v infer          1.0.7      v workflows      1.1.4
v modeldata      1.4.0      v workflowsets 1.1.0
v parsnip        1.2.1      v yardstick     1.3.1
v recipes        1.1.0

-- Conflicts ----- tidymodels_conflicts() --
x yardstick::accuracy() masks forecast::accuracy()
x data.table::between() masks dplyr::between()
x scales::discard() masks purrr::discard()
x Matrix::expand() masks tidyr::expand()
x rstan::extract() masks tidyr::extract()
x plotly::filter() masks dplyr::filter(), stats::filter()
x data.table::first() masks dplyr::first()
x recipes::fixed() masks stringr::fixed()
x jsonlite::flatten() masks purrr::flatten()
x kableExtra::group_rows() masks dplyr::group_rows()
x dplyr::lag() masks stats::lag()
x data.table::last() masks dplyr::last()
x caret::lift() masks purrr::lift()
x dials::mixture() masks brms::mixture()
x Matrix::pack() masks tidyr::pack()
x rsample::populate() masks Rcpp::populate()
x yardstick::precision() masks caret::precision()
x yardstick::recall() masks caret::recall()
x car::recode() masks dplyr::recode()
x MASS::select() masks plotly::select(), dplyr::select()
x yardstick::sensitivity() masks caret::sensitivity()
x xgboost::slice() masks plotly::slice(), dplyr::slice()
x car::some() masks purrr::some()
x yardstick::spec() masks readr::spec()
x yardstick::specificity() masks caret::specificity()
x recipes::step() masks stats::step()
x data.table::transpose() masks purrr::transpose()
x Matrix::unpack() masks tidyr::unpack()
x recipes::update() masks Matrix::update(), stats::update()
* Search for functions across packages at https://www.tidymodels.org/find/

[1] "All required libraries are loaded!"

Registered S3 method overwritten by 'GGally':
  method from
+.gg ggplot2

Attaching package: 'dataMaid'

The following object is masked from 'package:recipes':

  check

```


The following object is masked from 'package:infer':

visualize

The following object is masked from 'package:rmarkdown':

render

The following object is masked from 'package:lme4':

isSingular

The following object is masked from 'package:dplyr':

summarize

Rows: 23977 Columns: 6

-- Column specification -----

Delimiter: ","

chr (1): CountryName

dbl (4): Year, Value of global merchandise exports as a share of GDP, GDP pe...

lgl (1): World regions according to OWID

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
colnames(Trade)
```

```
[1] "CountryName"
[2] "Year"
[3] "Value of global merchandise exports as a share of GDP"
[4] "GDP per capita"
[5] "Population (historical)"
[6] "World regions according to OWID"
```

```
print(summary(Trade))
```

CountryName	Year
Length:23977	Min. : 1
Class :character	1st Qu.:1877
Mode :character	Median :1961
	Mean :1894
	3rd Qu.:1993
	Max. :2022

Value of global merchandise exports as a share of GDP	GDP per capita
Min. : 0.00	Min. : 295
1st Qu.: 11.16	1st Qu.: 1470
Median : 19.03	Median : 2618
Mean : 25.59	Mean : 6870
3rd Qu.: 31.36	3rd Qu.: 7232
Max. : 779.76	Max. : 160051

```

NA's      :10169                                NA's      :2391
Population (historical) World regions according to OWID
Min.      :8.821e+03      Mode:logical
1st Qu.:1.991e+06      NA's:23977
Median   :5.704e+06
Mean     :5.799e+07
3rd Qu.:1.921e+07
Max.     :8.021e+09
NA's     :4953

```

```
colnames(CW_Data)
```

```

[1] "CountryName"      "Year"
[3] "CurrentYearTotalTradeValue" "PreviousYearTotalTradeValue"
[5] "GrowthPercentage"  "AverageGrowthPercentage"

```

```
print(summary(CW_Data))
```

```

CountryName      Year      CurrentYearTotalTradeValue
Length:4554      Min.    :1989      Min.    :3.000e+00
Class :character  1st Qu.:2000      1st Qu.:1.634e+06
Mode  :character  Median :2007      Median :1.361e+07
                        Mean   :2007      Mean   :1.623e+08
                        3rd Qu.:2014      3rd Qu.:1.072e+08
                        Max.    :2021      Max.    :5.181e+09

PreviousYearTotalTradeValue GrowthPercentage  AverageGrowthPercentage
Min.    :3.000e+00          Min.    : -89.230      Min.    : -34.43
1st Qu.:1.528e+06          1st Qu.:  -4.685      1st Qu.:   5.81
Median :1.274e+07          Median :    6.830      Median :   8.31
Mean   :1.544e+08          Mean   :   12.361      Mean   :  12.36
3rd Qu.:9.970e+07          3rd Qu.:   18.538      3rd Qu.:  11.73
Max.   :4.997e+09          Max.   :5511.730      Max.   : 478.50

```

```
AnnVal <- read.csv("C:/Users/vasud/OneDrive/Desktop/U-M - ALL/STATS 506/Final Project Data/KAGGLE/archi
```

```

merged_data_full <- left_join(Trade, CW_Data, by = c("CountryName" = "CountryName", "Year" = "Year"))
merged_data_full2 <- left_join(merged_data_full, AnnVal, by = c("CountryName" = "CountryName", "Year" =

```

```
colnames(merged_data_full2)
```

```

[1] "CountryName"
[2] "Year"
[3] "Value of global merchandise exports as a share of GDP"
[4] "GDP per capita"
[5] "Population (historical)"
[6] "World regions according to OWID"
[7] "CurrentYearTotalTradeValue"
[8] "PreviousYearTotalTradeValue"
[9] "GrowthPercentage"
[10] "AverageGrowthPercentage"
[11] "AnnualTradeValue"

```

```
merged_data_full2 <- merged_data_full2[, !(names(merged_data_full2) %in% c("World regions according to (
colnames(merged_data_full2)
```

```
[1] "CountryName"
[2] "Year"
[3] "Value of global merchandise exports as a share of GDP"
[4] "GDP per capita"
[5] "Population (historical)"
[6] "CurrentYearTotalTradeValue"
[7] "PreviousYearTotalTradeValue"
[8] "GrowthPercentage"
[9] "AverageGrowthPercentage"
[10] "AnnualTradeValue"
```

```
TradeData <- merged_data_full2
colnames(TradeData)
```

```
[1] "CountryName"
[2] "Year"
[3] "Value of global merchandise exports as a share of GDP"
[4] "GDP per capita"
[5] "Population (historical)"
[6] "CurrentYearTotalTradeValue"
[7] "PreviousYearTotalTradeValue"
[8] "GrowthPercentage"
[9] "AverageGrowthPercentage"
[10] "AnnualTradeValue"
```

```
summary(TradeData)
```

CountryName	Year
Length:23977	Min. : 1
Class :character	1st Qu.:1877
Mode :character	Median :1961
	Mean :1894
	3rd Qu.:1993
	Max. :2022

Value of global merchandise exports as a share of GDP	GDP per capita
Min. : 0.00	Min. : 295
1st Qu.: 11.16	1st Qu.: 1470
Median : 19.03	Median : 2618
Mean : 25.59	Mean : 6870
3rd Qu.: 31.36	3rd Qu.: 7232
Max. : 779.76	Max. : 160051
NA's : 10169	NA's : 2391

Population (historical)	CurrentYearTotalTradeValue	PreviousYearTotalTradeValue
Min. : 8.821e+03	Min. : 1.250e+02	Min. : 1.950e+02
1st Qu.: 1.991e+06	1st Qu.: 2.674e+06	1st Qu.: 2.480e+06

```

Median :5.704e+06      Median :1.708e+07      Median :1.587e+07
Mean   :5.799e+07      Mean   :1.571e+08      Mean   :1.497e+08
3rd Qu.:1.921e+07      3rd Qu.:1.134e+08      3rd Qu.:1.072e+08
Max.   :8.021e+09      Max.   :5.181e+09      Max.   :4.997e+09
NA's   :4953           NA's   :20197          NA's   :20197

GrowthPercentage AverageGrowthPercentage AnnualTradeValue
Min.   : -86.12      Min.   : -34.43        Min.   :1.250e+02
1st Qu.: -3.95       1st Qu.: 5.81          1st Qu.:2.329e+06
Median : 7.06        Median : 8.42          Median :1.465e+07
Mean   : 10.83       Mean   : 10.61          Mean   :1.490e+08
3rd Qu.: 18.46       3rd Qu.: 11.57         3rd Qu.:1.031e+08
Max.   :1951.04      Max.   :292.12         Max.   :5.181e+09
NA's   :20197        NA's   :20197          NA's   :19947

```

```
summary(TradeData$AnnualTradeValue)
```

```

      Min.    1st Qu.    Median      Mean   3rd Qu.      Max.      NA's
1.250e+02 2.329e+06 1.465e+07 1.490e+08 1.031e+08 5.181e+09 19947

```

```
TradeData_clean <- na.omit(TradeData)
```

Cleaning and preparing the data

```

# Remove rows with NA values in specified columns
TradeData <- TradeData %>%
  filter(!is.na(`GDP per capita`) &
         !is.na(`Value of global merchandise exports as a share of GDP`) &
         !is.na(GrowthPercentage))

```

```
summary(TradeData)
```

```

CountryName      Year
Length:2680      Min.   :1989
Class :character  1st Qu.:1998
Mode  :character  Median :2004
                        Mean   :2003
                        3rd Qu.:2009
                        Max.   :2014

Value of global merchandise exports as a share of GDP GDP per capita
Min.   : 1.441                      Min.   : 561
1st Qu.: 16.718                      1st Qu.: 4333
Median : 25.572                      Median : 10694
Mean   : 31.336                      Mean   : 16165
3rd Qu.: 40.247                      3rd Qu.: 24513
Max.   :183.905                      Max.   :157713

Population (historical) CurrentYearTotalTradeValue PreviousYearTotalTradeValue
Min.   :6.830e+04        Min.   :5.177e+03      Min.   :5.177e+03
1st Qu.:3.796e+06        1st Qu.:3.210e+06      1st Qu.:2.966e+06
Median :1.008e+07        Median :1.968e+07      Median :1.769e+07

```

Mean	:4.799e+07	Mean	:1.482e+08	Mean	:1.386e+08
3rd Qu.:	2.922e+07	3rd Qu.:	1.161e+08	3rd Qu.:	1.067e+08
Max.	:1.388e+09	Max.	:4.685e+09	Max.	:4.418e+09
GrowthPercentage		AverageGrowthPercentage		AnnualTradeValue	
Min.	: -84.300	Min.	:-34.43	Min.	:5.177e+03
1st Qu.:	-1.093	1st Qu.:	5.91	1st Qu.:	3.210e+06
Median	: 9.205	Median	: 8.50	Median	:1.968e+07
Mean	: 12.557	Mean	: 10.33	Mean	:1.482e+08
3rd Qu.:	20.270	3rd Qu.:	11.70	3rd Qu.:	1.161e+08
Max.	:1951.040	Max.	: 88.73	Max.	:4.685e+09

```
library(car)

# Compute VIF
vif(lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP`, data = TradeData))

# GDP per capita
1.103906

# Value of global merchandise exports as a share of GDP
1.103906

# Check correlations between predictors
cor(TradeData[, c("GDP per capita", "Value of global merchandise exports as a share of GDP", "AnnualTradeValue")])
```

	GDP per capita	
GDP per capita	1.0000000	
Value of global merchandise exports as a share of GDP	0.3067997	
AnnualTradeValue	0.3809786	
	Value of global merchandise exports as a share of GDP	
GDP per capita	0.3067997	
Value of global merchandise exports as a share of GDP	1.0000000	
AnnualTradeValue	0.0438999	
	AnnualTradeValue	
GDP per capita	0.3809786	
Value of global merchandise exports as a share of GDP	0.0438999	
AnnualTradeValue	1.0000000	

1. **GDP per capita and Annual Trade Value:** The correlation is **0.38**, indicating a moderate positive relationship. This suggests that as GDP per capita increases, the annual trade value tends to increase as well, though not strongly.
2. **Value of Global Merchandise Exports as a Share of GDP and Annual Trade Value:** The correlation is quite **low** at **0.04**, indicating almost no linear relationship between these two variables. This implies that the share of merchandise exports as a portion of GDP is not strongly related to the total annual trade value.
3. **GDP per capita and Value of Global Merchandise Exports as a Share of GDP:** The correlation here is **0.31**, suggesting a weak to moderate positive relationship. This could indicate that richer countries (higher GDP per capita) tend to have a somewhat higher share of exports relative to GDP, but the relationship is not particularly strong.

```
# Run model with all predictors
lm_full <- lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP` + `AnnualTradeValue`)

# Compute VIF
vif(lm_full)
```

```
              `GDP per capita`
              1.297775
`Value of global merchandise exports as a share of GDP`
              1.111552
AnnualTradeValue
              1.177891
```

```
# Fit a reduced model (with three variables for VIF)
lm_reduced <- lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP` + `AnnualTradeValue`)

# Compute VIF
vif(lm_reduced)
```

```
              `GDP per capita`
              1.297775
`Value of global merchandise exports as a share of GDP`
              1.111552
AnnualTradeValue
              1.177891
```

```
library(mice)
```

Attaching package: 'mice'

The following object is masked from 'package:stats':

```
filter
```

The following objects are masked from 'package:base':

```
cbind, rbind
```

```
imputed_data <- mice(TradeData, m = 1, method = 'pmm', maxit = 10, seed = 123)
```

```
iter imp variable
```

```
1 1
2 1
3 1
4 1
5 1
6 1
7 1
8 1
```

```
9 1
10 1
```

Warning: Number of logged events: 2

```
TradeData_imputed <- complete(imputed_data)

# Perform PCA including the new variable
pca <- prcomp(TradeData_imputed[, c("GDP per capita", "Value of global merchandise exports as a share of GDP")])

# Get the first principal component
TradeData_imputed$PCA1 <- pca$x[, 1]

# Use the first principal component in the regression model
model_pca <- lm(GrowthPercentage ~ PCA1, data = TradeData_imputed)
summary(model_pca)
```

Call:

```
lm(formula = GrowthPercentage ~ PCA1, data = TradeData_imputed)
```

Residuals:

Min	1Q	Median	3Q	Max
-97.77	-13.66	-3.26	7.78	1938.23

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.5574	0.9813	12.797	<2e-16 ***
PCA1	-0.8157	0.7984	-1.022	0.307

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 50.8 on 2678 degrees of freedom

Multiple R-squared: 0.0003896, Adjusted R-squared: 1.635e-05

F-statistic: 1.044 on 1 and 2678 DF, p-value: 0.307

```
# Run the regression again with the new variable included
model <- lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP` + `AnnualTradeValue`, data = TradeData_imputed)
summary(model)
```

Call:

```
lm(formula = GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP` + `AnnualTradeValue`, data = TradeData_imputed)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.41	-13.62	-3.18	7.89	1938.90

Coefficients:

	Estimate	Std. Error
(Intercept)	1.213e+01	1.787e+00
`GDP per capita`	-1.058e-04	6.852e-05
`Value of global merchandise exports as a share of GDP`	7.814e-02	4.616e-02
AnnualTradeValue	-2.124e-09	2.903e-09

	t value	Pr(> t)
(Intercept)	6.789	1.38e-11 ***
`GDP per capita`	-1.545	0.1225
`Value of global merchandise exports as a share of GDP`	1.693	0.0906 .
AnnualTradeValue	-0.732	0.4644

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 50.77 on 2676 degrees of freedom
Multiple R-squared: 0.00221, Adjusted R-squared: 0.001091
F-statistic: 1.975 on 3 and 2676 DF, p-value: 0.1155

```
library(glmnet)

# Prepare data with the new column
x <- as.matrix(TradeData[, c("GDP per capita", "Value of global merchandise exports as a share of GDP",
y <- TradeData$GrowthPercentage

# Ridge regression (alpha = 0)
ridge_model <- glmnet(x, y, alpha = 0)

# Lasso regression (alpha = 1)
lasso_model <- glmnet(x, y, alpha = 1)

# Cross-validation to choose lambda
cv_ridge <- cv.glmnet(x, y, alpha = 0)
cv_lasso <- cv.glmnet(x, y, alpha = 1)

# Best lambda values
cv_ridge$lambda.min
```

```
[1] 11.77226
```

```
cv_lasso$lambda.min
```

```
[1] 0.00890528
```

```
TradeData_scaled <- TradeData
TradeData_scaled[, c("GDP per capita", "Value of global merchandise exports as a share of GDP", "Annual
coef(ridge_model, s = cv_ridge$lambda.min)
```

```
4 x 1 sparse Matrix of class "dgCMatrix"
```

	s1
(Intercept)	1.231165e+01

GDP per capita	-7.984363e-05
Value of global merchandise exports as a share of GDP	5.867669e-02
AnnualTradeValue	-2.039838e-09

```
coef(lasso_model, s = cv_lasso$lambda.min)
```

4 x 1 sparse Matrix of class "dgCMatrix"

	s1
(Intercept)	1.213900e+01
GDP per capita	-1.051822e-04
Value of global merchandise exports as a share of GDP	7.758722e-02
AnnualTradeValue	-2.109717e-09

Lambda values for Ridge and Lasso:

- Ridge (**lambda.min**): 11.77226
- Lasso (**lambda.min**): 0.00890528 These values represent the best regularization parameters selected through cross-validation.

1. Ridge regression coefficients at the optimal lambda (**lambda.min**):

- Intercept: 12.31165
- GDP per capita: -7.984363e-05
- Value of global merchandise exports as a share of GDP: 0.05867669
- AnnualTradeValue: -2.039838e-09

2. Lasso regression coefficients at the optimal lambda (**lambda.min**):

- Intercept: 12.13900
- GDP per capita: -1.051822e-04
- Value of global merchandise exports as a share of GDP: 0.07758722
- AnnualTradeValue: -2.109717e-09

Interpretation:

- Both models suggest a negative relationship between **GDP per capita** and the dependent variable (**GrowthPercentage**), and a positive relationship for the share of global merchandise exports and growth percentage.
- The coefficients for **AnnualTradeValue** in both models are very close to zero, suggesting that **AnnualTradeValue** has a minimal impact on the predicted **GrowthPercentage** after regularization

At this point, I want to check if adding the Annual Trade Value helped my model:

```
# Run a full linear model with all predictors including the new column
model_full <- lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP` + AnnualTradeValue)

# Summary of the model to check adjusted R-squared and other diagnostics
summary(model_full)
```

Call:

```
lm(formula = GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP`  
    AnnualTradeValue, data = TradeData)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.41	-13.62	-3.18	7.89	1938.90

Coefficients:

	Estimate	Std. Error
(Intercept)	1.213e+01	1.787e+00
`GDP per capita`	-1.058e-04	6.852e-05
`Value of global merchandise exports as a share of GDP`	7.814e-02	4.616e-02
AnnualTradeValue	-2.124e-09	2.903e-09

	t value	Pr(> t)
(Intercept)	6.789	1.38e-11 ***
`GDP per capita`	-1.545	0.1225
`Value of global merchandise exports as a share of GDP`	1.693	0.0906 .
AnnualTradeValue	-0.732	0.4644


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 50.77 on 2676 degrees of freedom

Multiple R-squared: 0.00221, Adjusted R-squared: 0.001091

F-statistic: 1.975 on 3 and 2676 DF, p-value: 0.1155

```
# Plot residuals to assess if the new predictor improves the model  
ggplot(TradeData, aes(x = fitted(model_full), y = residuals(model_full))) +  
  geom_point(alpha = 0.5) +  
  geom_hline(yintercept = 0, color = "red") +  
  labs(  
    title = "Residual Plot",  
    x = "Fitted Values",  
    y = "Residuals"  
  )
```



```
# Compare models by adjusted R-squared
model_no_new <- lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP`)
summary(model_no_new)$adj.r.squared
```

```
[1] 0.001264414
```

```
summary(model_full)$adj.r.squared
```

```
[1] 0.001091052
```

```
# Check for improvements in the model
if (summary(model_full)$adj.r.squared > summary(model_no_new)$adj.r.squared) {
  print("The new predictor improved the model.")
} else {
  print("The new predictor did not improve the model.")
}
```

```
[1] "The new predictor did not improve the model."
```

The new predictor did not necessarily improve the model.

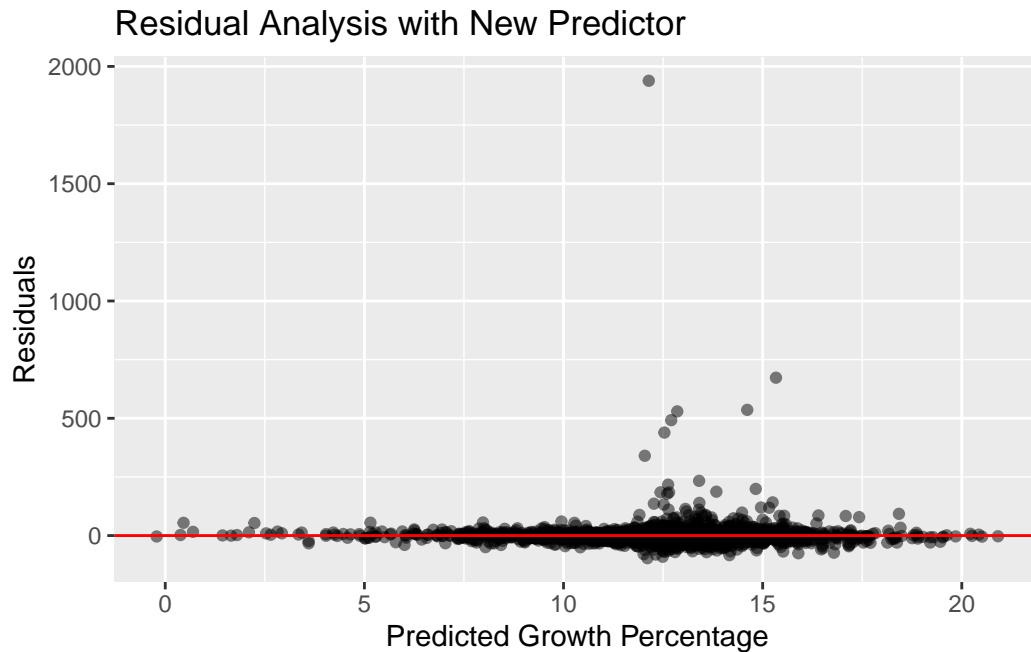
We can now do some predictive model analysis

```
# Run a linear model with all predictors including the new column
linear_model_full <- lm(GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP`)

# Add predictions to dataset
TradeData <- TradeData %>%
  mutate(PredictedGrowth = predict(linear_model_full, newdata = TradeData))

# Residual Analysis
ggplot(TradeData, aes(x = PredictedGrowth, y = GrowthPercentage - PredictedGrowth)) +
  geom_point(alpha = 0.5) +
```

```
geom_hline(yintercept = 0, color = "red") +
labs(
  title = "Residual Analysis with New Predictor",
  x = "Predicted Growth Percentage",
  y = "Residuals"
)
```



```
model_1 <- model_no_new
model_2 <- model_full

# Model Summaries
summary(model_1)
```

Call:

```
lm(formula = GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP`,
   data = TradeData)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.22	-13.60	-3.25	7.90	1939.19

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.205e+01	1.783e+00	6.756	1.74e-11 ***
`GDP per capita`	-1.252e-04	6.319e-05		
`Value of global merchandise exports as a share of GDP`	8.094e-02	4.599e-02		

```
`GDP per capita` -1.982 0.0476 *
`Value of global merchandise exports as a share of GDP` 1.760 0.0785 .
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 50.77 on 2677 degrees of freedom
```

```
Multiple R-squared:  0.00201,   Adjusted R-squared:  0.001264
```

```
F-statistic: 2.696 on 2 and 2677 DF,  p-value: 0.06767
```

```
summary(model_2)
```

```
Call:
```

```
lm(formula = GrowthPercentage ~ `GDP per capita` + `Value of global merchandise exports as a share of GDP`,
    data = TradeData)
```

```
Residuals:
```

```
      Min       1Q   Median       3Q      Max
-96.41  -13.62   -3.18    7.89 1938.90
```

```
Coefficients:
```

```

              Estimate Std. Error
(Intercept)    1.213e+01  1.787e+00
`GDP per capita` -1.058e-04  6.852e-05
`Value of global merchandise exports as a share of GDP`  7.814e-02  4.616e-02
AnnualTradeValue -2.124e-09  2.903e-09

              t value Pr(>|t|)
(Intercept)    6.789 1.38e-11 ***
`GDP per capita` -1.545  0.1225
`Value of global merchandise exports as a share of GDP`  1.693  0.0906 .
AnnualTradeValue -0.732  0.4644
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 50.77 on 2676 degrees of freedom
```

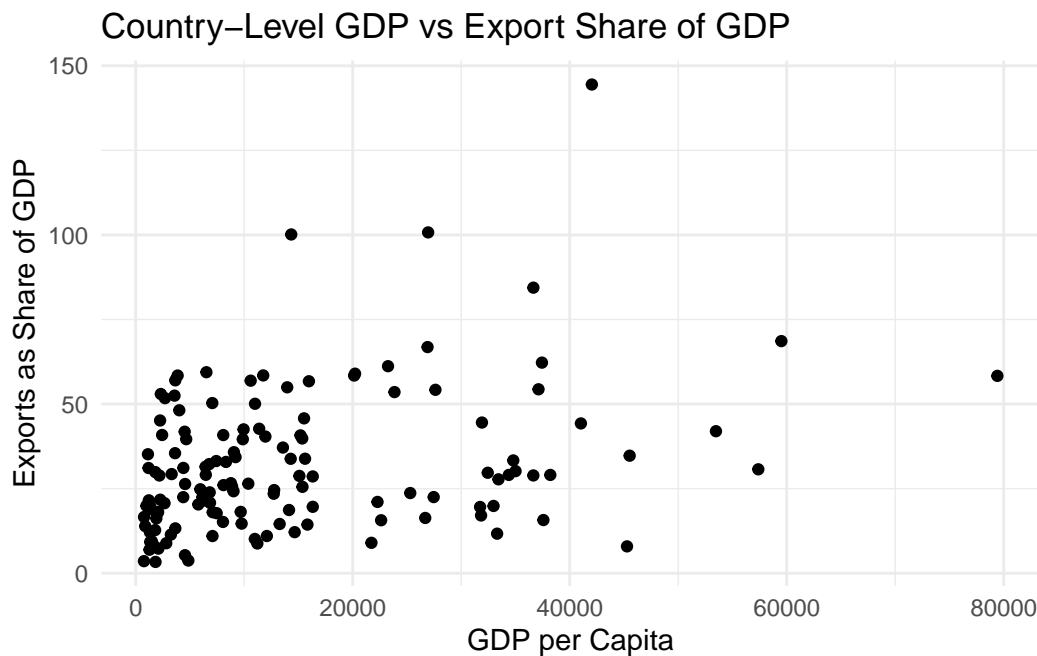
```
Multiple R-squared:  0.00221,   Adjusted R-squared:  0.001091
```

```
F-statistic: 1.975 on 3 and 2676 DF,  p-value: 0.1155
```

```
TradeData_country_summary <- TradeData %>%
  group_by(CountryName) %>%
  summarise(mean_gdp_per_capita = mean(`GDP per capita`, na.rm = TRUE),
            mean_exports_share_of_GDP = mean(`Value of global merchandise exports as a share of GDP`, na.rm = TRUE))

# EDA Visualizations
# Scatterplot: GDP per Capita vs. Exports
# Exporting model summaries
broom::tidy(model_1) %>% write.csv("Model_1_Summary.csv")
broom::tidy(model_2) %>% write.csv("Model_2_Summary.csv")
```

```
# Visualize GDP per capita vs exports share for different countries
ggplot(TradeData_country_summary, aes(x = mean_gdp_per_capita, y = mean_exports_share_of_GDP)) +
  geom_point() +
  labs(title = "Country-Level GDP vs Export Share of GDP",
        x = "GDP per Capita",
        y = "Exports as Share of GDP") +
  theme_minimal()
```



```
# Further statistical analysis: e.g., correlation between GDP per capita and export share
cor(TradeData_clean$`GDP per capita`,
     TradeData_clean$`Value of global merchandise exports as a share of GDP`,
     use = "complete.obs")
```

```
[1] 0.3067997
```

```
# For time series analysis, group by Year and summarize export share
TradeData_year_summary <- TradeData_clean %>%
  group_by(Year) %>%
  summarise(mean_exports_share_of_GDP = mean(`Value of global merchandise exports as a share of GDP`, na.rm = TRUE))

# Plot export share over the years
ggplot(TradeData_year_summary, aes(x = Year, y = mean_exports_share_of_GDP)) +
  geom_line(color = "red") +
  labs(title = "Global Merchandise Exports as Share of GDP Over Time",
        x = "Year",
        y = "Mean Share of GDP") +
  theme_minimal()
```



```
# Saving cleaned data for further analysis or export
write_csv(TradeData_clean, "cleaned_data.csv")
```

Both models suggest that the predictors (**GDP per capita**, **Value of global merchandise exports as a share of GDP**, and **AnnualTradeValue**) explain very little of the variation in growth percentage, as indicated by the low R-squared values. Only the **GDP per capita** variable in the first model is statistically significant at the 5% level, with a negative relationship to growth percentage. The export share of GDP has a marginal significance (near the 10% level) but does not show a strong effect.

Thus, the models indicate weak explanatory power and suggest that other factors not included in the model may have a stronger influence on growth percentage.

Advanced Analysis

```
# Interaction terms in regression
model_3 <- lm(`Value of global merchandise exports as a share of GDP` ~
              `GDP per capita` * GrowthPercentage, data = TradeData)
summary(model_3)
```

Call:

```
lm(formula = `Value of global merchandise exports as a share of GDP` ~
    `GDP per capita` * GrowthPercentage, data = TradeData)
```

Residuals:

Min	1Q	Median	3Q	Max
-41.887	-14.014	-5.579	9.059	138.749

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.453e+01	6.207e-01	39.514	<2e-16 ***

```

`GDP per capita`          4.097e-04  2.804e-05  14.614  <2e-16 ***
GrowthPercentage          7.380e-05  1.549e-02   0.005    0.996
`GDP per capita`:GrowthPercentage 1.037e-06  9.633e-07   1.077    0.282
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 21.32 on 2676 degrees of freedom
Multiple R-squared: 0.09556, Adjusted R-squared: 0.09455
F-statistic: 94.25 on 3 and 2676 DF, p-value: < 2.2e-16

```

# Exporting interaction model summary
broom::tidy(model_3) %>% write.csv("Model_3_Summary.csv")

# Feature Engineering
TradeData <- TradeData %>%
  mutate(
    Population_in_millions = TradeData$"Population (historical)" / 1e6,
    GDP_to_population_ratio = `GDP per capita` / TradeData$"Population (historical)"
  )

```

```

# View the summarized country-level data
head(TradeData_country_summary)

```

```

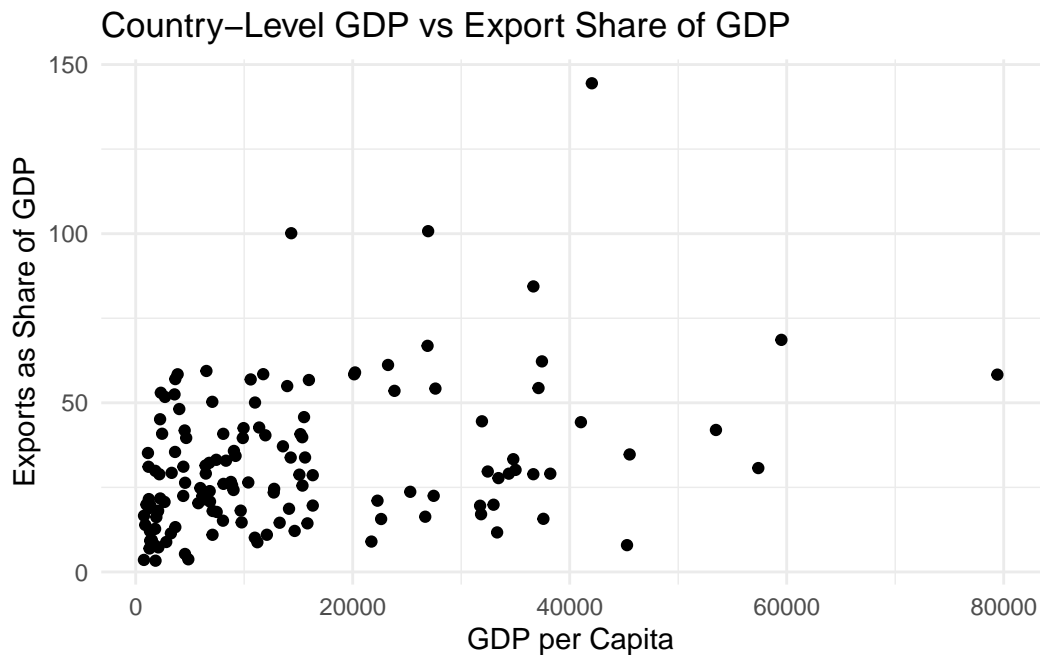
# A tibble: 6 x 3
  CountryName mean_gdp_per_capita mean_exports_share_of_GDP
  <chr>          <dbl>          <dbl>
1 Afghanistan    1828.           3.33
2 Albania         7080.          11.0
3 Algeria         9055.          35.8
4 Angola          7060.          50.3
5 Argentina      15829.          14.4
6 Armenia         8044.          15.2

```

```

# Visualize GDP per capita vs exports share for different countries
ggplot(TradeData_country_summary, aes(x = mean_gdp_per_capita, y = mean_exports_share_of_GDP)) +
  geom_point() +
  labs(title = "Country-Level GDP vs Export Share of GDP",
       x = "GDP per Capita",
       y = "Exports as Share of GDP") +
  theme_minimal()

```

```
# Further statistical analysis: e.g., correlation between GDP per capita and export share
cor(TradeData_clean$`GDP per capita`,
    TradeData_clean$`Value of global merchandise exports as a share of GDP`,
    use = "complete.obs")
```

```
[1] 0.3067997
```

```
# For time series analysis, group by Year and summarize export share
TradeData_year_summary <- TradeData_clean %>%
  group_by(Year) %>%
  summarise(mean_exports_share_of_GDP = mean(`Value of global merchandise exports as a share of GDP`, na.rm = TRUE))

# Plot export share over the years
ggplot(TradeData_year_summary, aes(x = Year, y = mean_exports_share_of_GDP)) +
  geom_line(color = "red") +
  labs(title = "Global Merchandise Exports as Share of GDP Over Time",
       x = "Year",
       y = "Mean Share of GDP") +
  theme_minimal()
```



```
# Saving cleaned data for further analysis or export  
write_csv(TradeData_clean, "cleaned_data.csv")
```